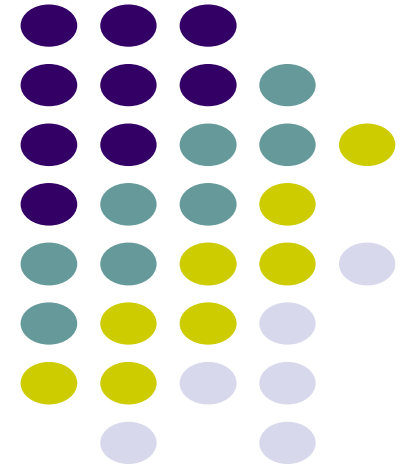



CS 528: Mobile and Ubiquitous Computing

Lecture 5b: Facial Analysis: How it works

Emmanuel Agu







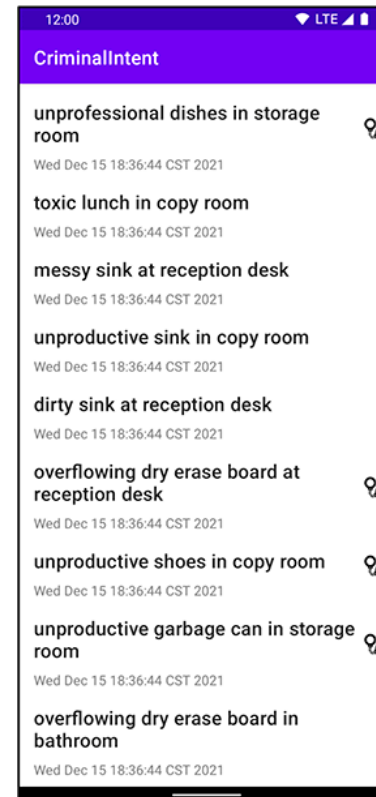
Criminal Intent

ANR Skipped Chapters

Chapter 13: Fragment Navigation



- Fragment navigation used to get list and details parts of **CriminalIntent** working together
- Uses Navigation jetpack library
- When user presses item in list of crimes:
 - Navigation library swaps out **CrimeListFragment**
 - Replaces it with new instance of **CrimeDetailFragment**





Chapter 16: Implicit Intents

- In **CriminalIntent**:
 - Enable picking a suspect for a crime from user's contact list
 - Use implicit intent to send a text-based report of a crime

A screenshot of a mobile application titled "CriminalIntent". The interface includes a status bar at the top showing the time 12:00 and battery level. Below the title bar, there is a "Title" section with a text input field containing the placeholder "Enter a title for the crime.". This is followed by a "Details" section containing a date and time string "SAT JAN 22 13:17:26 CST 2022" and a checkbox labeled "Solved". At the bottom of the form, there are two prominent blue buttons: "CHOOSE SUSPECT" and "SEND CRIME REPORT".

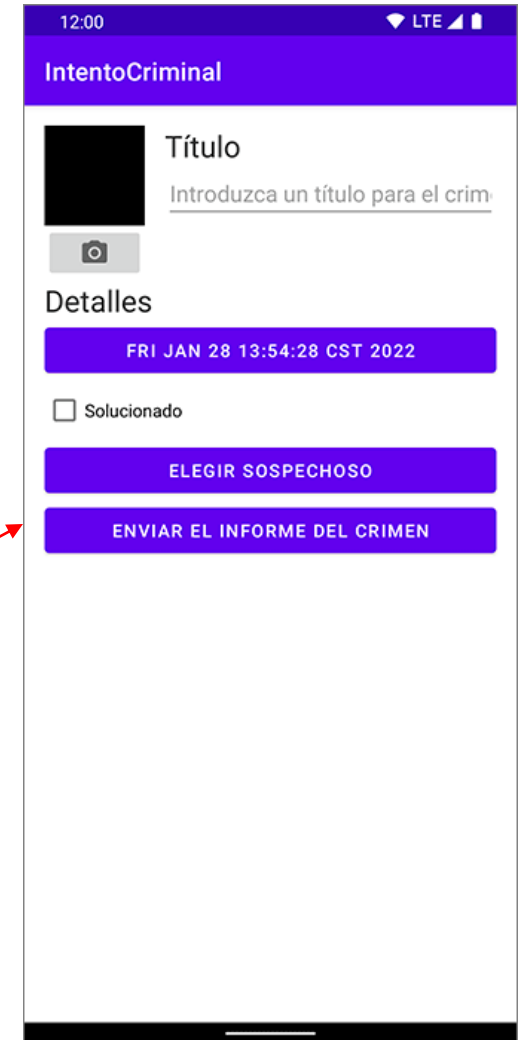
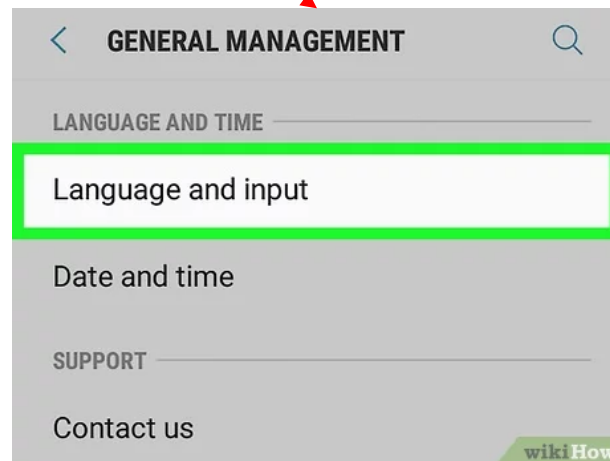
Open contacts app

Open an app that knows how to send text



Chapter 18: Localization

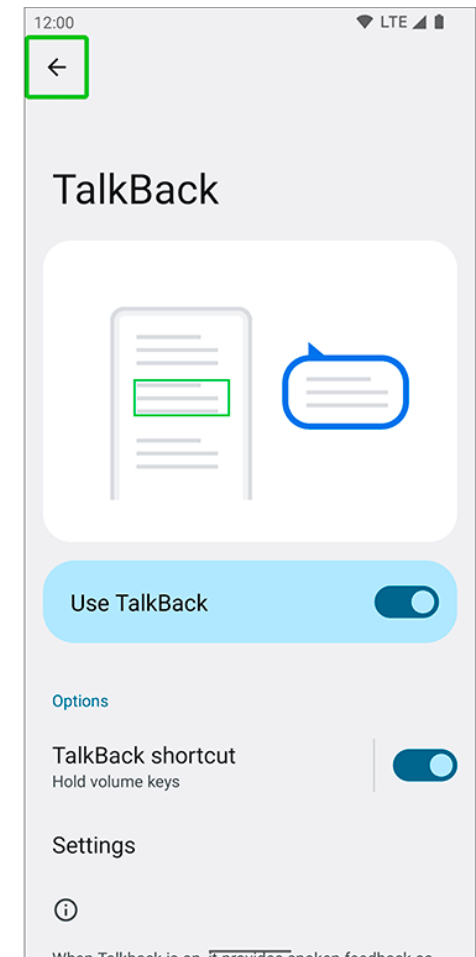
- **Localization:** provide app's resources in various languages
- In **CriminalIntent**, localize text that user sees so that app can be read in Spanish or English
 - Create a Spanish version of strings.xml
 - When device language is set to Spanish, android runtime uses Spanish strings.




Chapter 19: Accessibility

- **Accessibility:** Makes app usable by people with impaired vision, mobility or hearing
- **Goal:** Make **CriminalIntent** more accessible for people with impaired vision
 - By reading text out loud using Talkback
- **Talkback:** Android screen reader made by Google, speaks out screen's content depending on what user is doing
- E.g. when app opened and + new crime clicked, Talkback speaks out "New Crime. Double-tap to activate"

Navigate Up button.
Double-tap to activate.





Final Project



Final Project: 3 main flavors of projects

Project Requirements

For the final project, you are required to propose a project that might:

1. Design and develop an Android app that solves a real problem that WPI students have. The app is to use mobile or ubiquitous computing components (e.g. location, sensors or camera) and your projects difficulty will be graded based on the difficulty points sheet [[HERE](#)] . You can also enrich your application by pulling data/information from various third party/web sources.
2. Research and develop an app that classifies some human sensor data . E.g.
 - Classifies a speaker's voice to determine if nervous, sad, etc
 - Classifies a person's accelerometer data and detects them from 5-10 other people
 - Classifies a picture of a person's face and determines their mood
 - Classifies data from a person's phone to measure their sleep duration or/and quality
 - Detects a person's heart rate from a video of their face
 - Classifies person's communication/phone usage patterns to detect their mood
3. Programs an Android Deep/Machine Learning module (e.g. TensorFlow)



Final Project: Timeline

Deliverable	Description	Deadline
1-Slide:	Submit 1 slide describing your final project idea	Thursday, October 12, 2023 by class time
15-min proposal pitch	Give a 15-minute pitch of your project in class	Thursday, November 2, 2023, in class
Submit proposal pitch slides	Submit slides for your 15-minute pitch of your project in class	Friday, November 3, 2023, 11.59PM
Final project presentation:	Present your final project in class (15 mins)	Thursday, December 14, 2023, in class
Final paper submission:	Submit a final paper describing your project, your final project presentation slides and all your code	Friday, December 15, 2023, 11.59PM



Final Project: Difficulty Score

- **Project execution, presentation, paper: 80%**
- **Project difficulty score: 20%**
- **Mobile Components and Android UI (4 points each)**
 - Every 5 Android screens (A maximum of 8 points can be earned for the UI)
 - Playback audio/video
 - Maps, location sensing
 - Camera: taking pictures
- **Ubiquitous Computing Components & Android UI (6 points each)**
 - Activity Recognition, sensor programming, step counting
 - GeoFencing, ML Kit modules: e.g. Face/barcode detection/tracking
- **Machine/Deep Learning (10 points each)**
 - Machine/deep learning (i.e. run study, gather data or use existing dataset to classify/detect something)
 - Program Android, machine learning/deep learning components

Final Project: Generating Ideas/Past Projects



Generating Ideas

If you're having difficulties coming up with project ideas, check out the following links:

- [[Final Projects from CS 528 \(this class\) in S23](#)]
- [[Final Projects from CS 528 \(this class\) in F20](#)]
- [[Final Projects from CS 528 \(this class\) in F19](#)]
- [[Final Projects from CS 528 \(this class\) in F18](#)]
- [[Final Projects from CS 528 \(this class\) in S16](#)]
- [[Final Projects from CS 4518 \(undergrad class\) in C17](#)]
- [[Final Projects from CS 403X \(undergrad class\) in D16](#)]
- [[My collection of ubicomp project links](#)]
- [[Personal Informatics tools page](#)]


Final Project: 1-slide



- 1-slide from group due on Thursday, October 12
 - 2/30 of final project grade
- Slide should cover 3 aspects
 1. **Problem you intend to work on either:**
 - App that solves a real problem that WPI students face
 - Points awarded for difficulty, components used (location, sensor, camera, ML)
 - If games, must gamify solution to real world problem
 2. **Why this problem is important**
 - E.g. WPI students want to find study space close to end of term
 3. **Summary of envisioned mobile app (?) solution**
 1. E.g. Mobile app tracks availability of study spaces. WPI students can check using app
- You can:
 - Bounce ideas of me (email, or in person)
 - Change idea any time

Reminder: Detection vs Recognition

Ref: Adriana Kovashka, Introduction to Vision (CS 1674), U. of Pittsburgh, fall 2016



- Detection: Find (draw squares around) all faces in image
- Recognition: Whose face is in the image



Detection




Recognition

“Sally”



Uses of Face Detection

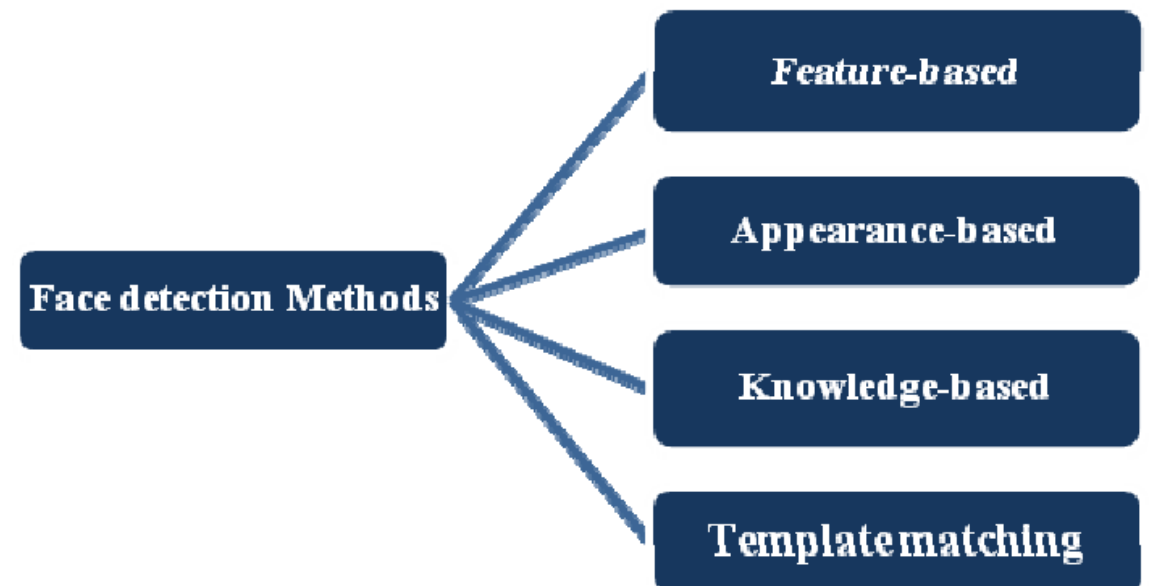
- Detect Faces in surveillance images
- Cameras use it to detect faces in a picture
- Facebook: detect faces in an image, recognize them





Types of Face Detection Algorithms

- Categorization by Yang, M.H., Kriegman, D.J. and Ahuja, N., 2002. Detecting faces in images: A survey. *IEEE Transactions on pattern analysis and machine intelligence*, 24(1), pp.34-58.
- **Knowledge-based:** Rules based on human knowledge. E.g. relative positions and distances of parts (nose, eyes, mouth, etc)
- **Feature-based:** Extract visual attributes (color, texture, shape), train machine learning classifier to distinguish facial vs. non-facial regions
- **Template:** Divide face into parts (e.g. mouth, eyes), compare/match parts to templates (or standard face patterns)
- **Appearance-based:** Generates templates from representative set of faces





Face Detection: Viola Jones Algorithm

- Most popular method, published in CVPR 2001 (top computer vision conference)
- Windows based: Draw candidate windows, decides if it contains face or not
- First fast, real time, still used today
- Challenge?
 - Lots of pixels in image (millions)
 - But faces are rare (0 – 10 per image)
 - Window evaluation must be computationally efficient, fast
 - But also low false positive rate (< 1 in 1 million)
 - False Positive: Say it's a face but it's not

ACCEPTED CONFERENCE ON COMPUTER VISION AND PATTERN RECOGNITION 2001

Rapid Object Detection using a Boosted Cascade of Simple Features

Paul Viola
viola@merl.com
Mitsubishi Electric Research Labs
201 Broadway, 8th FL
Cambridge, MA 02139

Michael Jones
mjones@crl.dec.com
Compaq CRL
One Cambridge Center
Cambridge, MA 02142


Abstract

This paper describes a machine learning approach for vi-

ected at 15 frames per second on a conventional 700 MHz Intel Pentium III. In other face detection systems, auxiliary information, such as image differences in video sequences,

Voila Jones: 4 Stages

1. Haar Feature Selection
2. Creating an Integral Image
3. Adaboost Training
4. Cascading Classifiers





Step 1: Haar Feature Selection



- Features: rectangular patterns
- Calculate sums of pixel values within rectangles

- Example:

- Overlay feature A over image
- Calculate (sum of pixels in white rectangle) – (sum of pixels in gray rectangle)




- So what?

- All faces have common properties
- E.g.
 - Eye region is darker than upper cheeks
 - Bridge of nose is brighter than eyes



Step 2: Integral Image


- Efficient way to calculate sum of pixel values within rectangles
 - Calculate (sum of pixels in white rectangle) – (sum of pixels in gray rectangle)






Step 3: AdaBoost Classifier

- Want to overlay various scales, and positions of the 4 feature types over the face
 - Determine which specific shapes, positions, scales are discriminative
- Lots of permutations and combinations: E.g. 180,000 features per 24 x 24 pixel window
- Use an AdaBoost classifier (Machine Learning)




Boosting: Training

Ref: Adriana Kovashka, Introduction to Vision (CS 1674), U. of Pittsburgh, fall 2016




- Initially, weight each training example equally
- In each boosting round:
 - a) Find the weak learner that achieves the lowest weighted training error
(Find dividing line that has lowest total distance from all the points in the training set)




Boosting: Training

Ref: Adriana Kovashka, Introduction to Vision (CS 1674), U. of Pittsburgh, fall 2016




- In each boosting round:
 - Find the weak learner that achieves the lowest weighted training error
 - Increase weights of training examples misclassified by current weak learner**
(Increase weights of examples on wrong side of dividing line)




Boosting: Training

Ref: Adriana Kovashka, Introduction to Vision (CS 1674), U. of Pittsburgh, fall 2016




- Round 2: Repeat a) and b) again to find weak classifier 2
 - a) **Find the weak learner that achieves the lowest weighted training error**
(Note: weighted examples results in different dividing line)




Boosting: Training

Ref: Adriana Kovashka, Introduction to Vision (CS 1674), U. of Pittsburgh, fall 2016




- Round 2: Repeat a) and b) again to find weak classifier 2
 - a) Find the weak learner that achieves the lowest weighted training error
 - b) **Increase weights of training examples misclassified by current weak learner**
(Increase weights of examples on wrong side of dividing line)

**Weights
Increased**




Boosting: Training

Ref: Adriana Kovashka, Introduction to Vision (CS 1674), U. of Pittsburgh, fall 2016




- Round 3: Repeat a) and b) again to find weak classifier 3




Boosting: Training

Ref: Adriana Kovashka, Introduction to Vision (CS 1674), U. of Pittsburgh, fall 2016



- Final classifier is combination of weak classifiers.
- i.e. for each point to be classified, test against:
 - Classifier 1
 - Classifier 2
 - Classifier 3... etc


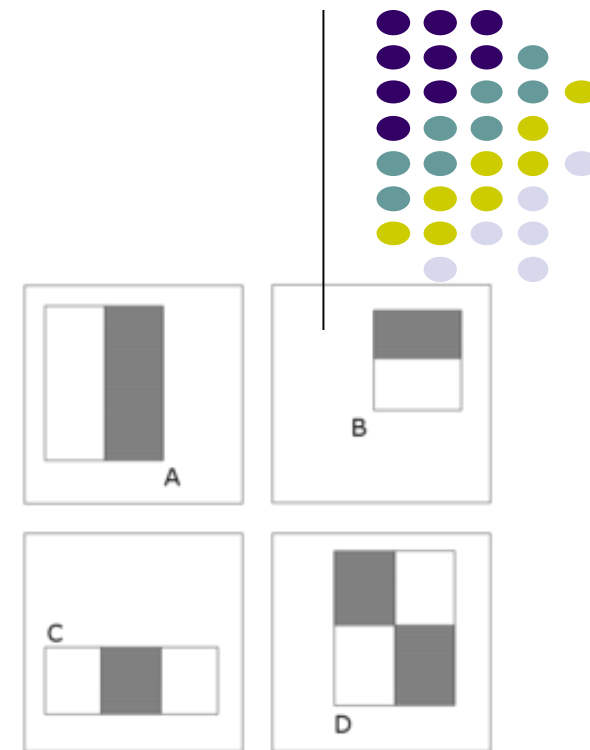
**Final classifier is
a combination of weak
classifiers**



- Different algorithms, formulas for re-weighting and combining weak learners
 - E.g. AdaBoost

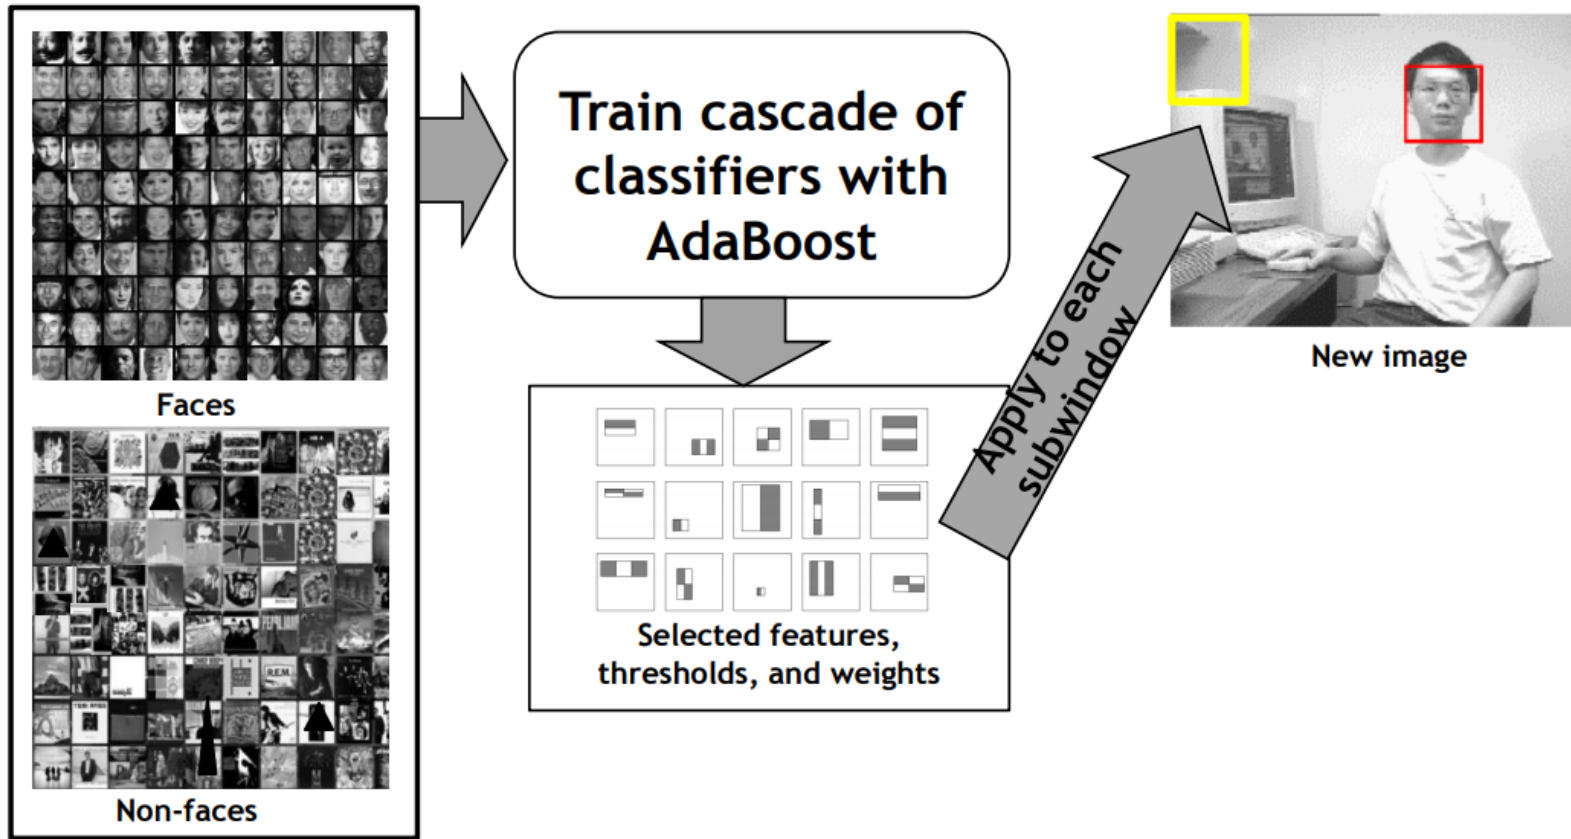
AdaBoost Classifier

- Examine windows within an image
- What to determine if there's a face in that window
- Overlay 4 feature patterns over various parts of the window
- Find which of 4 feature patterns + locations that indicate that window contains face
 - Patterns + position + scale that discriminate face vs. non-face used as weak classifiers
 - Use Boosted combination of multiple weak classifiers as final classifier
 - Form cascade of weak classifiers, reject negatives quickly





Voila Jones Algorithm: Summary



Training: slow





Test: fast

- Integral image
- Cascade of classifiers

Train with 5K positives, 350M negatives



Real-time detector using 38 layer cascade (0.067s)

Voila Jones: Results



Face Recognition: Uses

- Device security:
 - Recognize owner's face, use as their device password
- Identify genetic disorders:
 - Analyze, compare faces to databases of people with various disorders
 - E.g. Face2gene app, DeepGestalt software
- Prevent shoplifting
 - Recognize past shoplifters, notify owner if they visit
- Check underage alcohol/tobacco buyers
 - Automatically assess buyers age from their face
- Security in schools
 - Recognize criminals, sex offenders
- Airline ticketing:
 - Use your face as your ticket





Android Face Recognition/Face Unlocking

- Mostly used for security, recognize owners face in place of password
 - Not definitely more secure but faster, more convenient (no forgotten passwords)
- Basic version compares stored picture with picture captured by front-facing camera
- Thief can fool system by presenting a picture of owner
- Available on Pixel 4, well executed
 - Cannot be fooled by similar people, photos





Face Recognition: Triplet Loss

- Uses neural networks a lot these days
 - Ref: Schroff *et al*, 2015, FaceNet: A Unified embedding for face recognition and clustering

- Given:

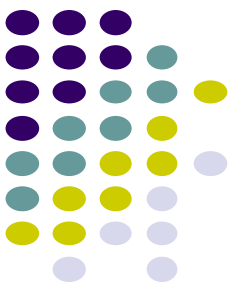
- Anchor: reference image
- Positive: Positive example
- Negative: Negative example

Triplet Loss

Learning Objective



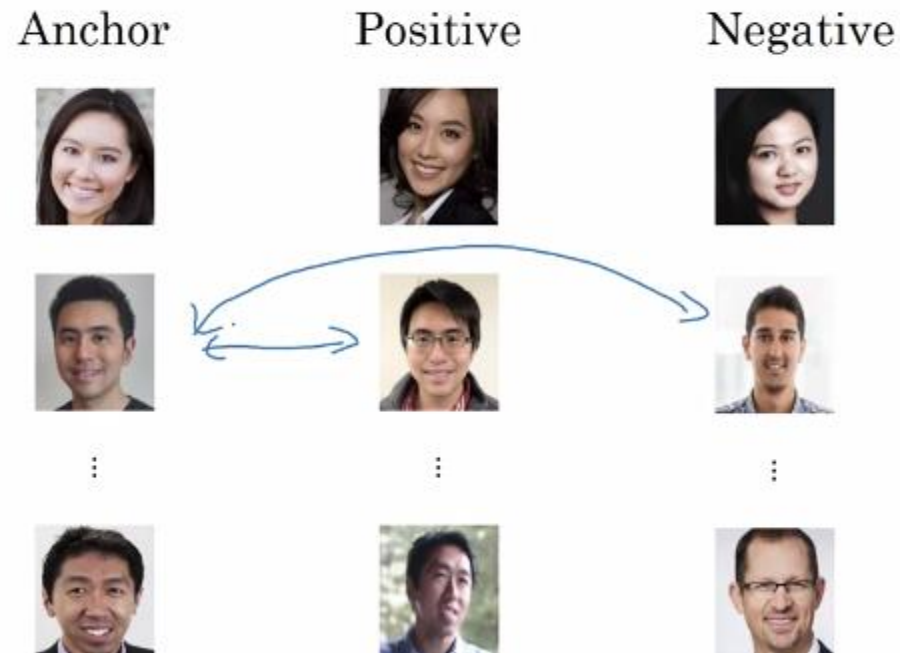
- Defines loss (objective function) that learns encoding (0101000...) such that:
 - Anchor and positive close together
 - Anchor and negative far apart

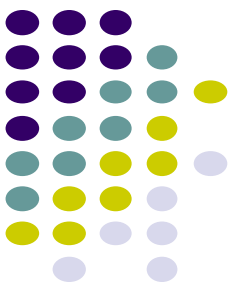


Triplet Loss: Training Set

- Millions of sets of 3 (Anchor, Positive, Negative)
- Neural network learns encoding and to distinguish faces

Training set using triplet loss





References

- Android Nerd Ranch, 5th edition
- Google Android Tutorials
- Adriana Kovashka, Introduction to Vision (CS 1674), U. of Pittsburgh, fall 2016
- Divyansh Dwivedi, Face Detection For Beginners,
<https://towardsdatascience.com/face-detection-for-beginners-e58e8f21aad9>
- The complete guide to Facial recognition, Panda Security, October 11, 2019,
<https://www.pandasecurity.com/mediacenter/panda-security/facial-recognition-technology/>
- Voila Jones Object Detection Framework Wikipedia page,
https://en.wikipedia.org/wiki/Viola%E2%80%93Jones_object_detection_framework

References



- Network programming with Android, <https://slideplayer.com/slide/8471042/>
- Android Networking, <https://developer.android.com/training/basics/network-ops>
- Convolutional Neural Networks Course, Deeplearning.ai, Coursera