


# Mobile and Ubiquitous Computing

## Lecture 8a: Final Project, Paper Presentations, Introduction to Machine Learning

**Emmanuel Agu**





# Final Project Proposal

# Final Project Proposal



- Final project 1 slide due, submitted on Tue (Oct 24, 10am)?
- While working on project 3, more brainstorming on final project
- November 2 (Next Thursday), Propose final project (mobile/ubicom app or machine learning project that solves a real WPI problem)
- All teams will present next week!!
- Proposals should include:
  1. **Problem you intend to work on**
    - App that finds available study spaces (safe + available), dynamically updated
  2. **Why this problem is important**
    - E.g. 32% of WPI students living with roommates, hard to find places to study
  3. **Related Work:** What prior solutions have been proposed for this problem

# Final Project Proposal



## 4. Summary of envisioned solution

- E.g. Mobile app maintains dynamic list of available and safe study spots including Android/third party modules app will have

## 5. Tally your difficulty points in your slides, summarize your tally

- Can also do Machine learning project that classifies/detects analyzes a dataset of builds a real-time app to classify some human sensor data.
  - Can use existing smartphone datasets online, or gather your own data
- You can:
  - Bounce ideas of me (email, or in person)
  - Change idea any time

# Rubric: Grading Considerations



- **Problem (10/100)**
  - How much is the problem a real problem (e.g. not contrived)
  - Is this really a good problem that is a good fit to solve with mobile/ubiquitous computing? (e.g. are there better approaches?)
  - How useful would it be if this problem is solved?
  - What is the potential impact on the community (e.g. WPI students) (e.g. how much money? Time? Productivity.. Would be saved?)
  - What is the evidence of the importance? (E.g. quote a statistic)

# Rubric: Grading Considerations




- **Related Work (5/100)**
  - Prior research, other apps been previously proposed to solve this problem?
- **Proposed Solution/Classification (10/100)**
  - How good/clever/interesting is the solution?
  - How sophisticated and how are the mobile/ubiquitous computing components (high level) used? (e.g. location, geofencing, activity recognition, face recognition, machine learning, etc)

# Rubric: Grading Considerations



- **Implementation Plan + Timeline (10/100)**
  - Clear plans to realize your design/methodology
  - Android modules/3<sup>rd</sup> party software used
  - Software architecture,
  - Screenshots (or sketches of UI), or study design + timeline
- **Evaluation Plan (5/100)**
  - How will you evaluate your project, metrics
  - E.g. small user studies for apps
  - Machine learning performance metrics (e.g. classification accuracy, F1 score, etc)
- **Difficulty Points (20/100)**
  - Will follow rubric handed out in class, and scale max. of 25 down to 20/100
- 40 more points allotted for your slides + oral presentation



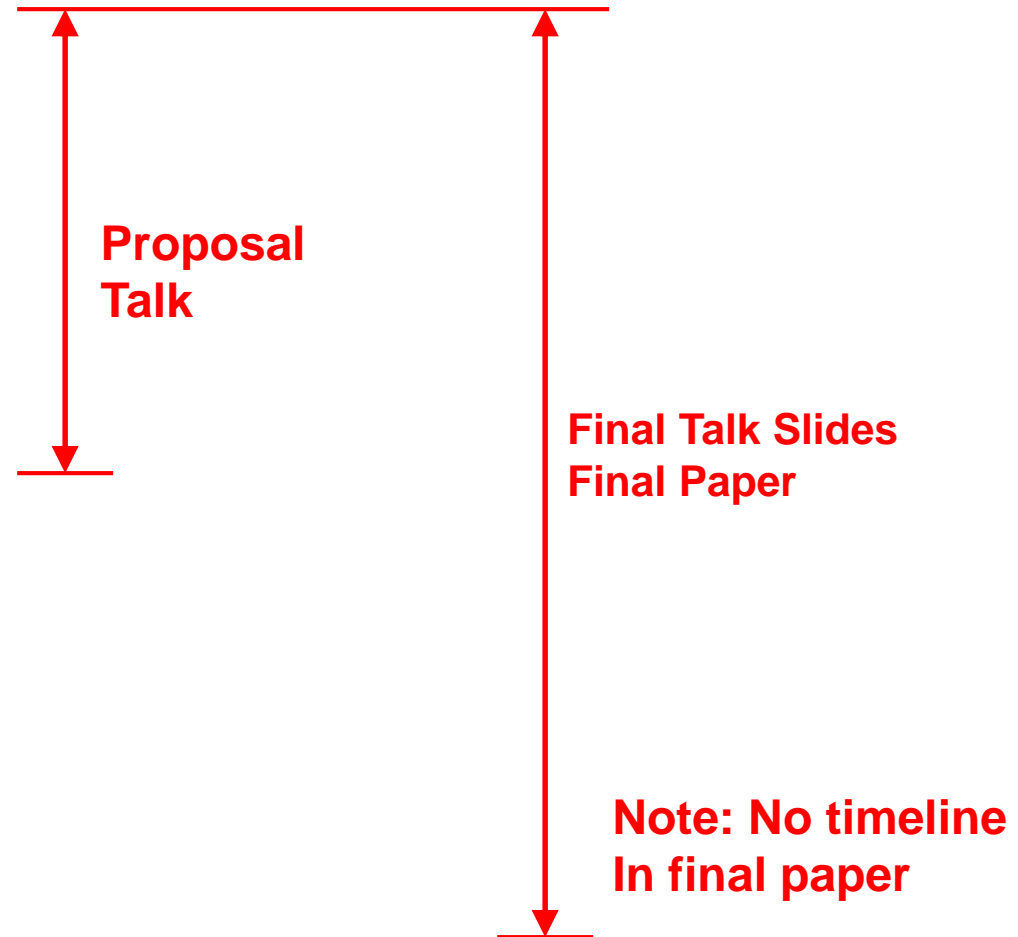
# **Final Project: Proposal Vs Final Submission (Presentation + Paper)**






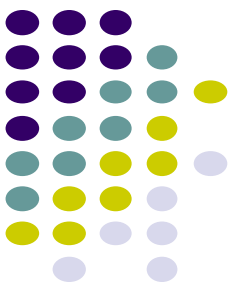
## Final Project Proposal Vs Final Submission

- Introduction
- Related Work
- Approach/methodology
- Implementation
- **Project timeline**
- Evaluation/Results
- Discussion
- Conclusion
- Future Work





# Student Research Paper Presentations



# Research Presentation: Mobile and UbiComp Papers

- I have put up list of research papers on Canvas
- On Nov. 16 and 30, GROUPs present 1 research paper each from my list
- Your talk should cover:
  - Motivation for problem (General)
  - Specific problem solved in paper
  - Approach used to solve the problem/how it works
  - Evaluation of solution (sample results)
  - Discussion/conclusion



# **Intuitive Introduction to Machine Learning for Ubiquitous Computing**





## My Goals in this Section

- If you know machine learning
  - Set off light bulb
  - Projects involving ML?
- If you don't know machine learning
  - Get general idea, how it's used
- Knowledge will also make papers easier to read/understand

# Recall: Activity Recognition

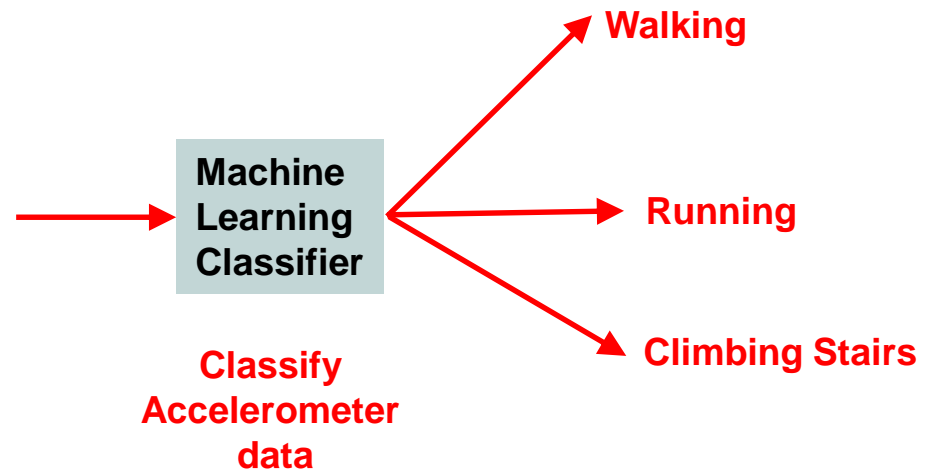
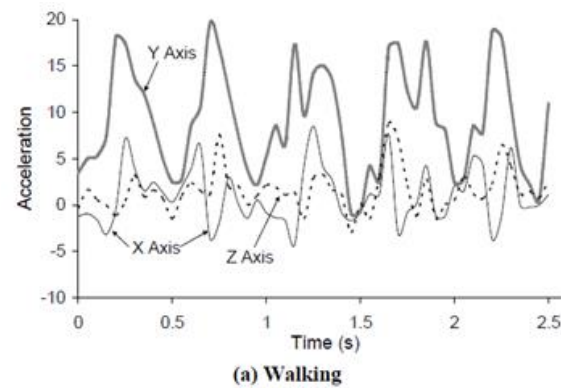
- Want app to detect when user is performing any of the following 6 activities
  - Walking,
  - Jogging,
  - Ascending stairs,
  - Descending stairs,
  - Sitting,
  - Standing



# Recall: Activity Recognition Overview




**Gather Accelerometer data**




# Recall: Example Accelerometer Data for Activities


Different user activities generate different accelerometer patterns




(a) Walking



(b) Jogging



(e) Sitting


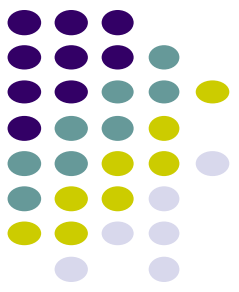


(f) Standing




# Recall: Example Accelerometer Data for Activities

Different user activities generate different accelerometer patterns




(c) Ascending Stairs

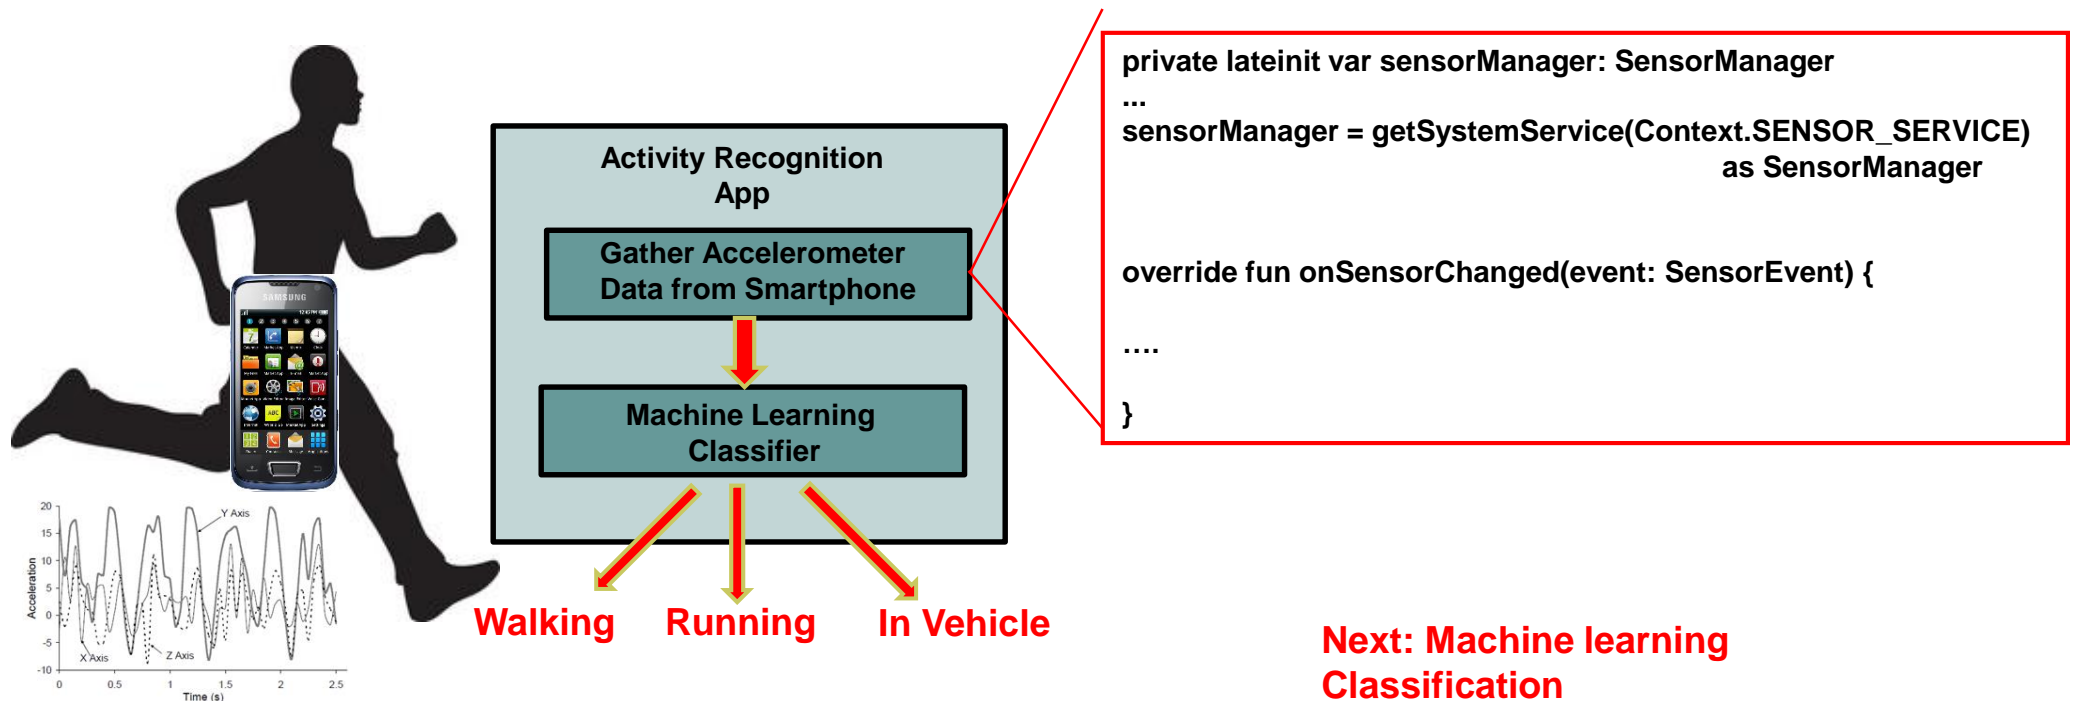


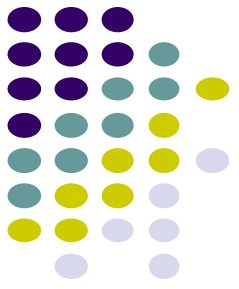
(d) Descending Stairs

# DIY Activity Recognition (AR) Android App




- As user performs an activity, AR app on user's smartphone
  - Gathers accelerometer data
  - Uses **machine learning classifier** to determine what activity (running, jumping, etc) accelerometer pattern corresponds to
- Classifier:** Machine learning algorithm that guesses what activity **class** (or type) accelerometer sample corresponds to



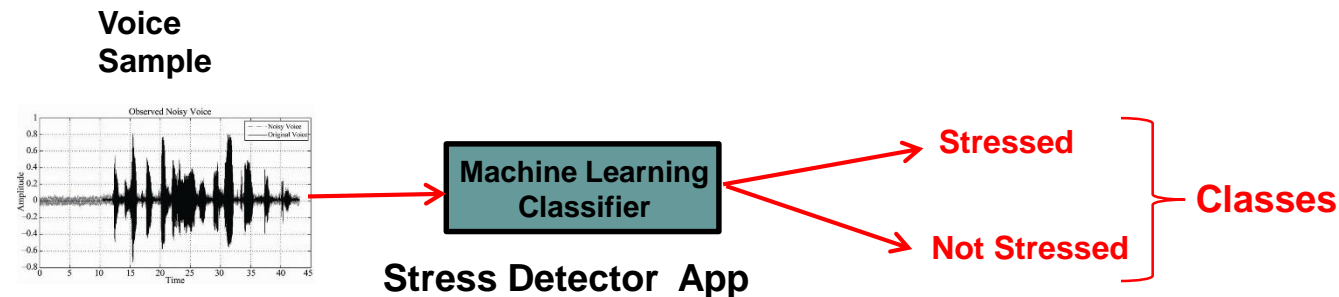
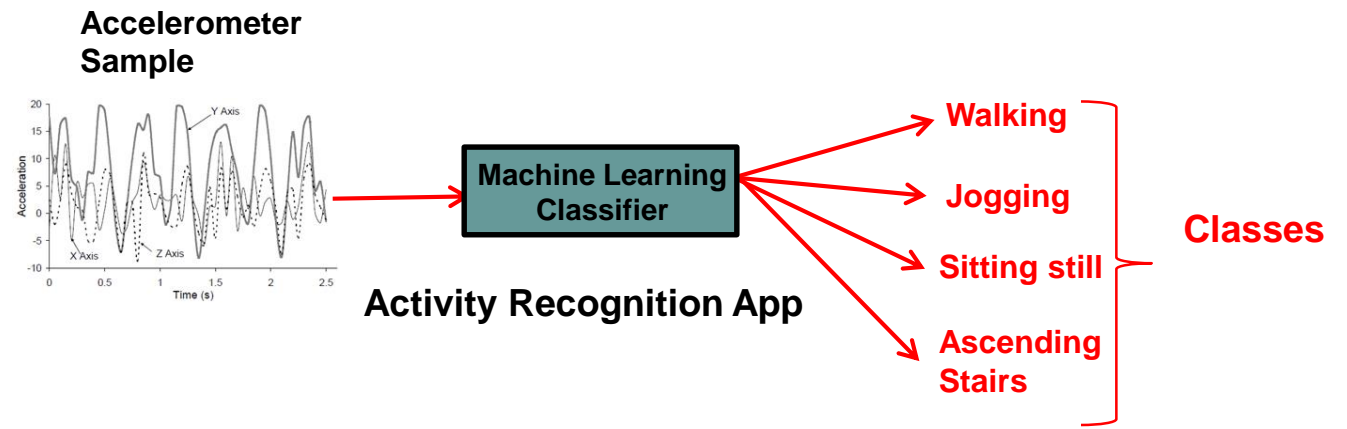


# Classification for Ubiquitous Computing



# Classification

- **Classification** is type of machine learning used a lot in Ubicomp
- Classification? determine which **class** a sample (e.g. snippet of accelerometer data) belongs to. Examples:



# Classification


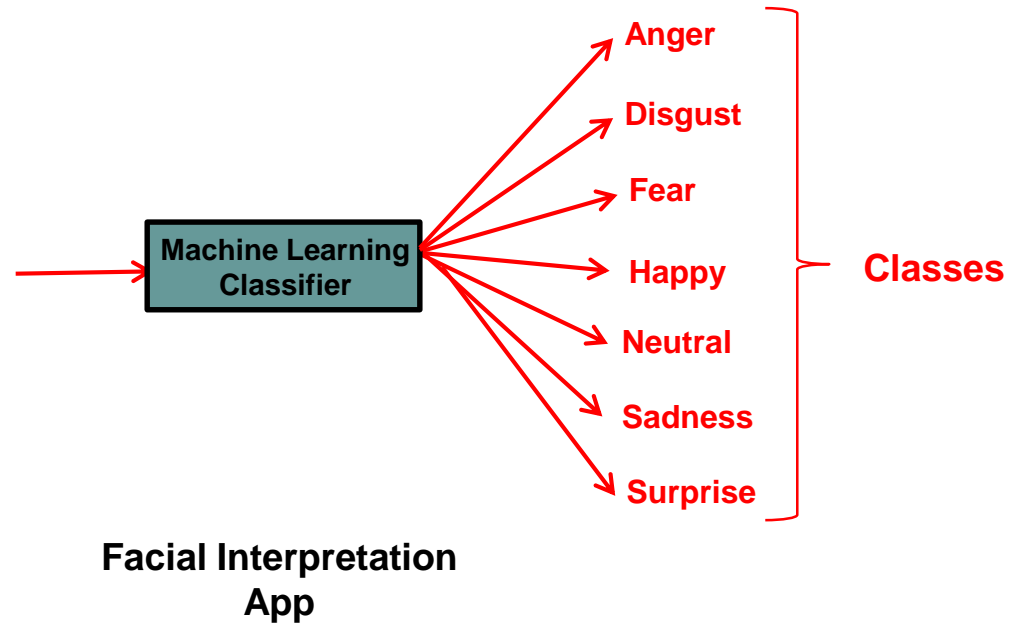



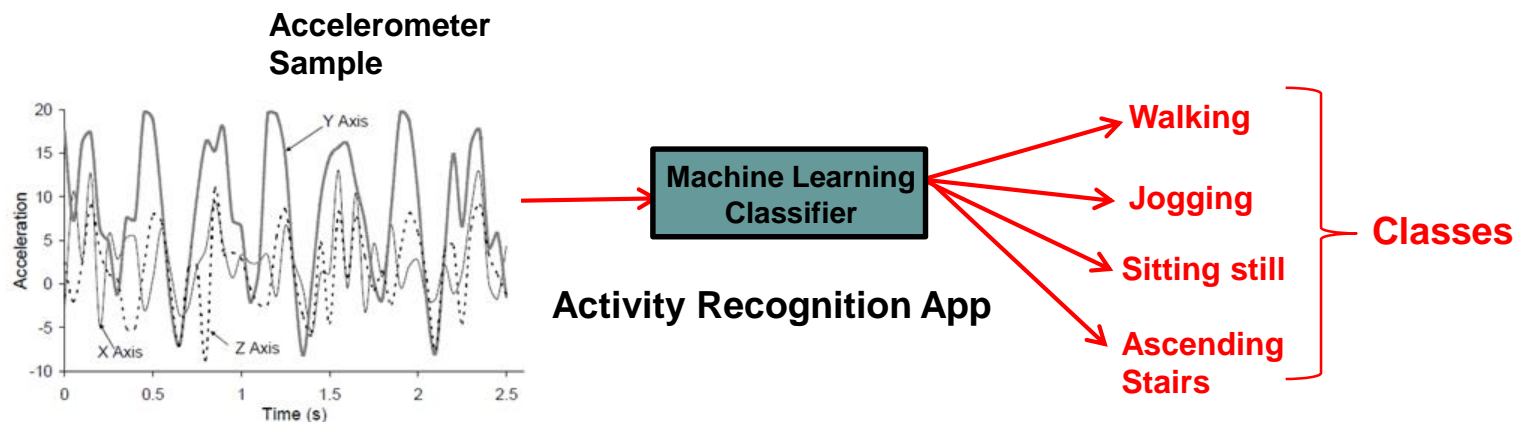
Image showing  
Facial Expression



# Classifier

- Analyzes new sample, guesses corresponding class
- Intuitively, can think of classifier as set of rules for classification. E.g.
- Example rules for classifying accelerometer signal in Activity Recognition

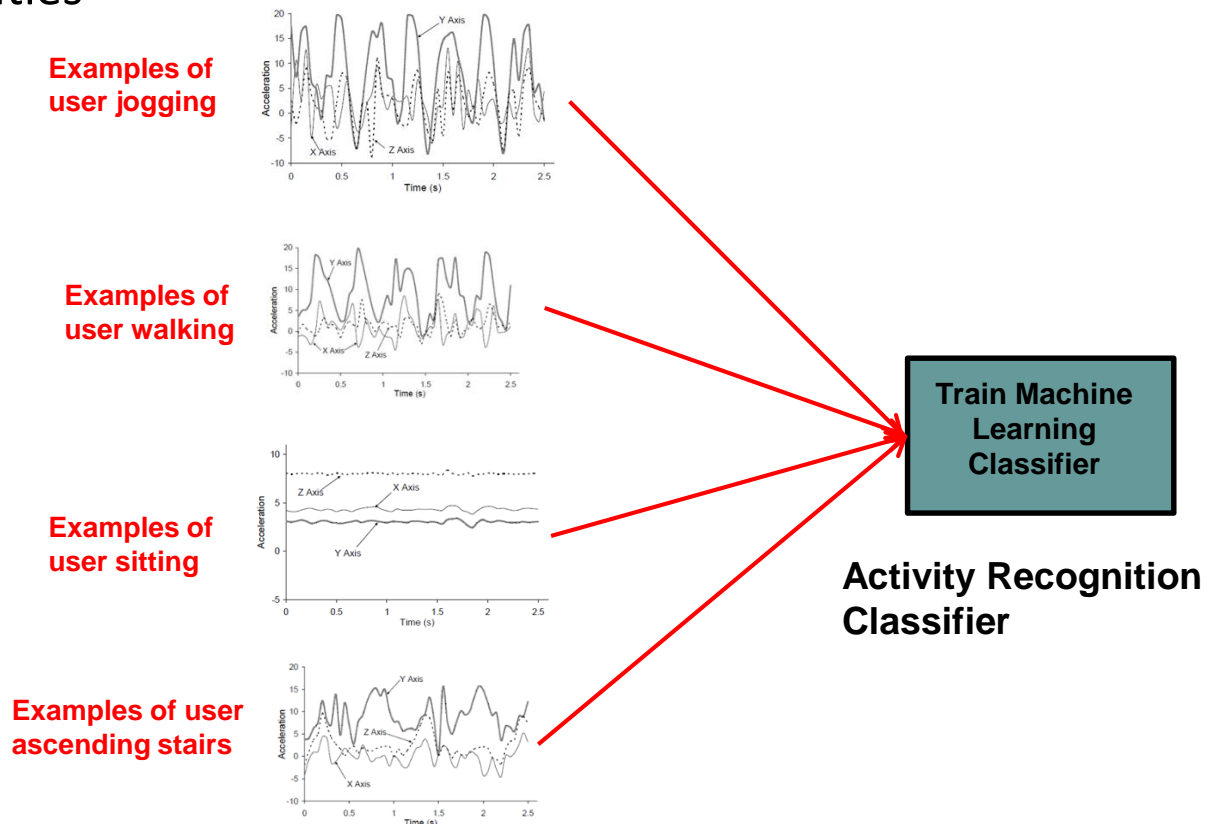
```
If ((Accelerometer peak value > 12 m/s)
    and (Accelerometer average value < 6 m/s)) {
    Activity = "Jogging";
}
```

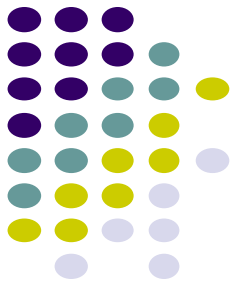




# Training a Classifier

- Created using example-based approach (called training)
- **Training a classifier:** Given examples of each class => generate rules (ML model) to categorize new samples
- **E.g:** Analyze 30+ Examples (from 30 subjects) of accelerometer signal for each activity type (walking, jogging, sitting, ascending stairs) => generate rules (classifier) to classify future activities





# Training a Classifier: Steps






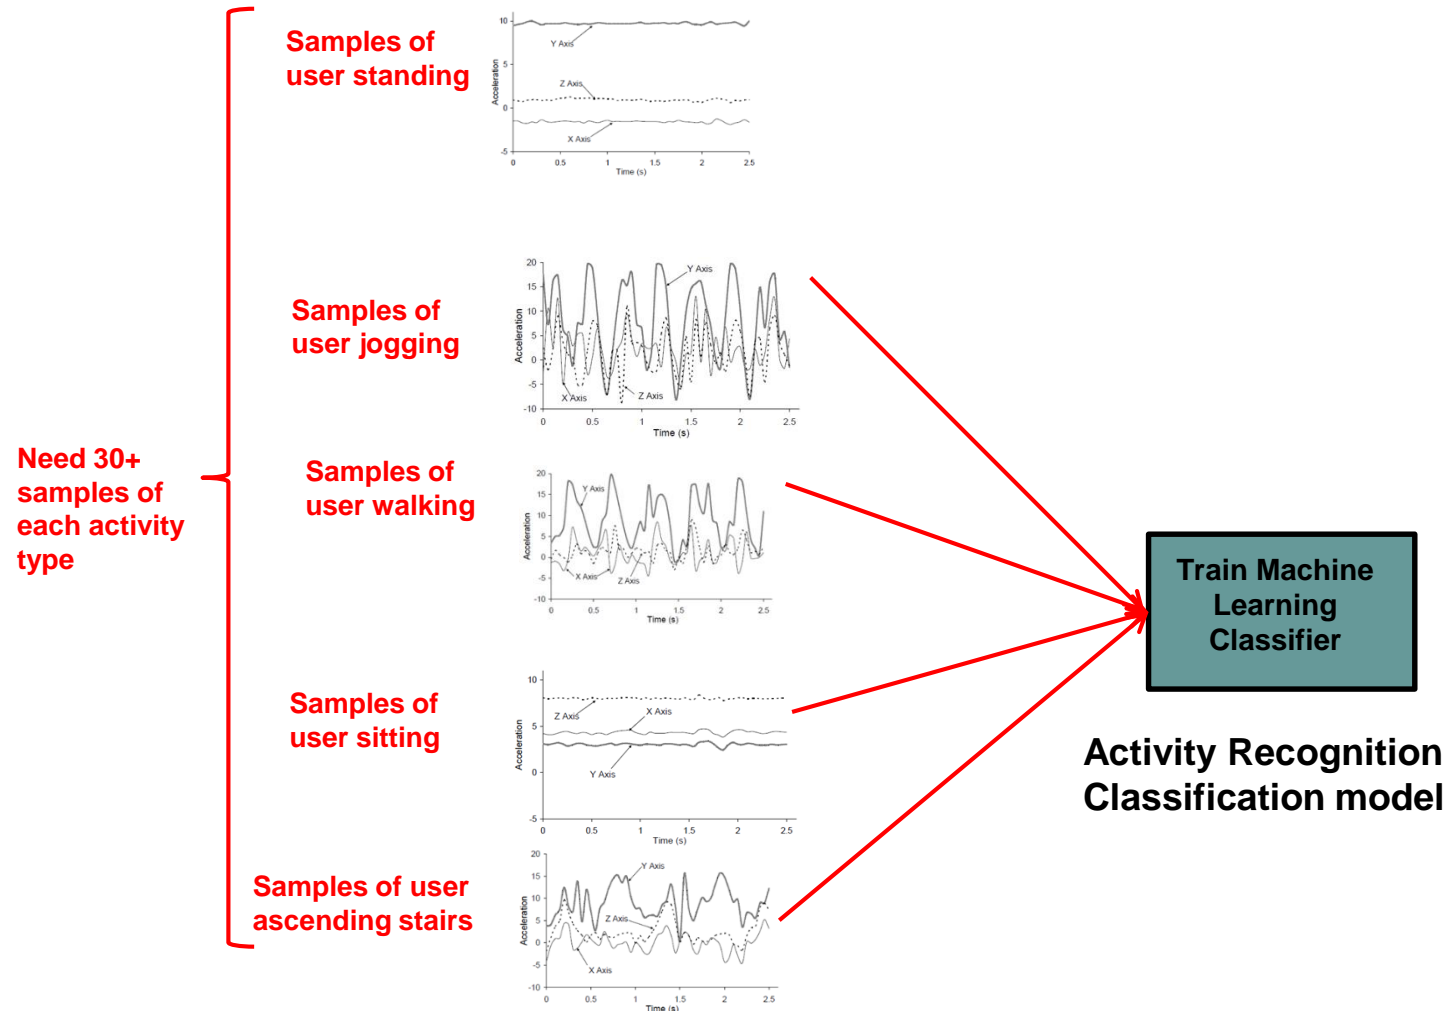
# Steps for Training a Classifier

1. Gather data samples + label them
2. Import accelerometer samples into classification library (e.g. scikit-learn, MATLAB)
3. Pre-processing (segmentation, smoothing, etc)
4. Extract features
5. Train classifier
6. Deploy classifier

# Step 1: Gather Sample data + Label them



- Need many samples of accelerometer data corresponding to each activity type (jogging, walking, sitting, ascending stairs, etc)





## Step 1: Gather Sample data + Label them

- Conduct a study to gather sample accelerometer data for each activity class
  - Recruit 30+ subjects
  - Run app that gathers accelerometer sensor data on subject's phone
  - Each subject:
    - Perform each activity (walking, jogging, sitting, etc)
    - Collect accelerometer data while they perform each activity (walking, jogging, sitting, etc)
  - Label data. i.e. tag each accelerometer sample with the corresponding activity
- Now have 30+ examples of each activity






Figure 1: Axes of Motion Relative to User

**30+  
Samples of  
user sitting**



**30+ Samples of  
user ascending  
stairs**





## Step 1: Gather Sample data + Label them

### Program to Gather Accelerometer Data


- **Option 1:** Can write “home-grown” sensor program app that gathers accelerometer data while user is doing each activity (1 at a time)

```
private lateinit var sensorManager: SensorManager
...
sensorManager = getSystemService(Context.SENSOR_SERVICE) as SensorManager

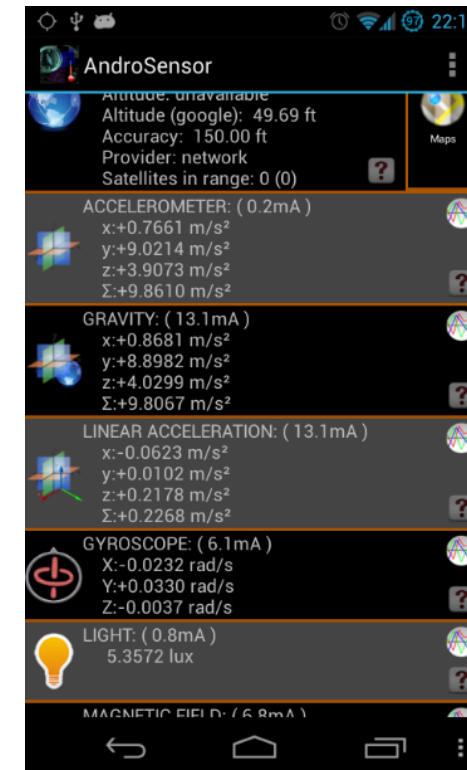
override fun onSensorChanged(event: SensorEvent) {
....
}
```

# Step 1: Gather Sample data + Label them

## Program to Gather Accelerometer Data

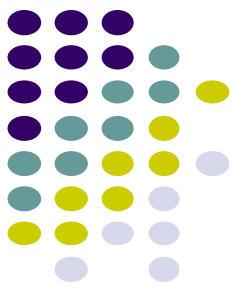


- **Option 2:** Use 3<sup>rd</sup> party app to gather accelerometer
  - E.g. **AndroSensor**
  - Just download app,
    - Select sensors to log (e.g. accelerometer)
    - Continuously gathers sensor data in background
    - Saves sensor data to .csv file

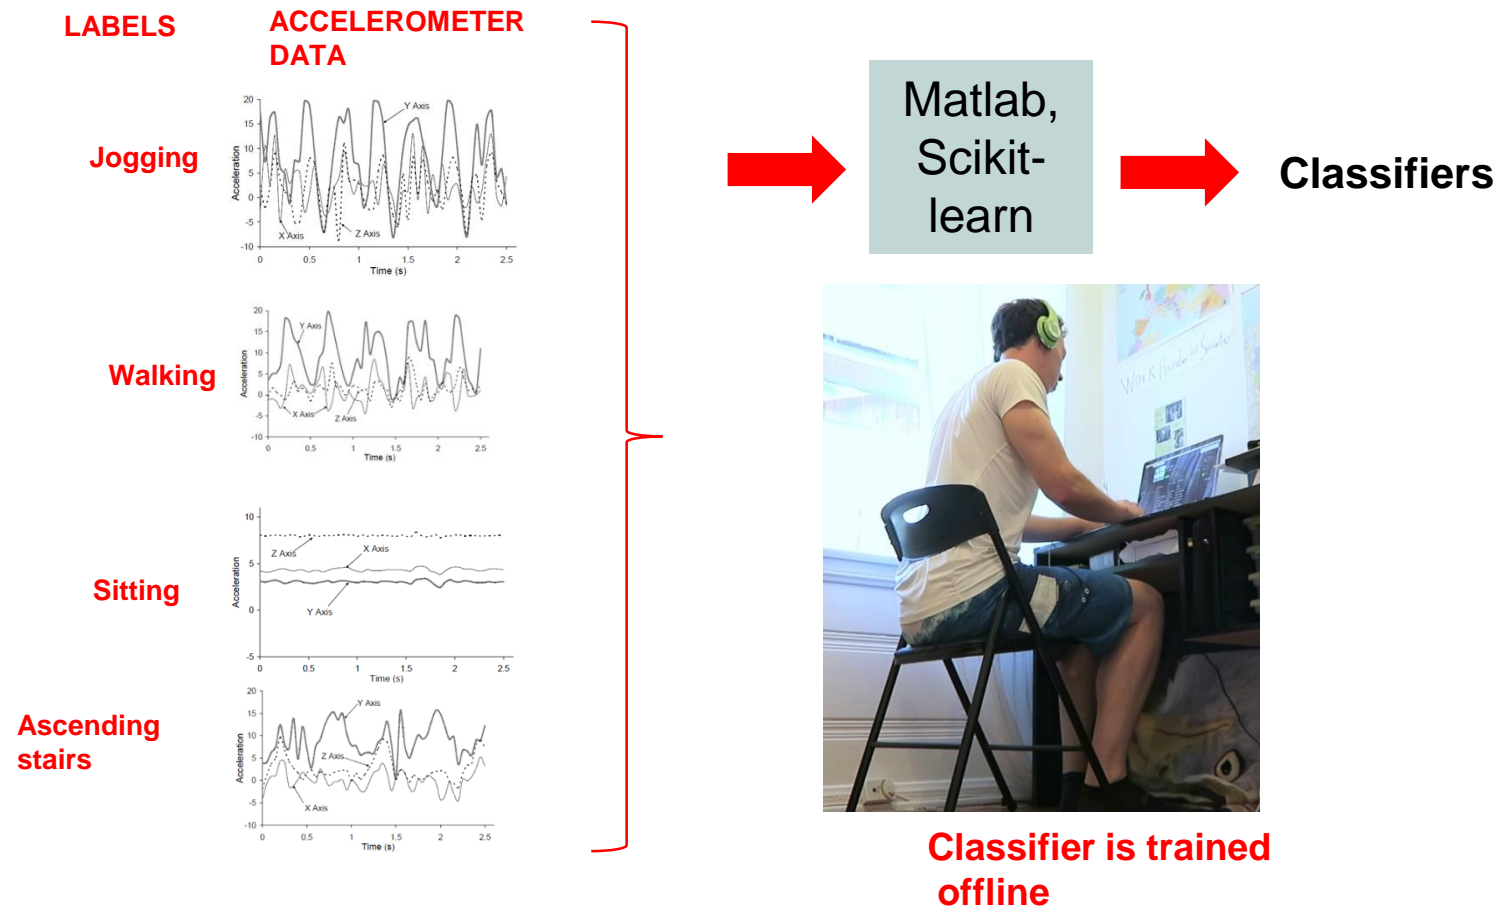


**AndroSensor**

## Step 2: Import accelerometer samples into classification library (e.g. Scikit-Learn, MATLAB)



- Import accelerometer data (labelled with corresponding activity) into MATLAB, scikit-learn (or other Machine learning Framework)







## Step 3: Pre-processing (segmentation, smoothing, etc)

### Segment Data (Windows)


- Pre-processing data (in scikit-learn or MATLAB) may include segmentation, smoothing, etc
  - **Segment:** Divide data into smaller chunks. E.g. divide 60 seconds of raw time-series data into 5 second chunks
    - Note: 5 seconds of accelerometer data may be 100s of readings
  - **Smoothing:** Replace groups of raw values with moving average





## Step 4: Compute (Extract) Features

- For each 5-second segment (batch of accelerometer values) compute features (in scikit-learn, MATLAB, etc.)
- **Features:** Formulas to quantify attributes, characteristics of accelerometer data
- **Example features calculated using data in each segment:**
  - Minimum value
  - Maximum value
  - min-max of values
  - Largest magnitude
  - Average
  - Standard deviation
  - ... various statistics



(e) Sitting







## Step 4: Compute (Extract) Features

- **Important:** Ideally, feature values are different for, can distinguish each activity type (class)
- **E.g:** Consider min-max range feature

Large min-max  
for jogging



Small min-max  
for sitting





## Step 4: Compute (Extract) Features

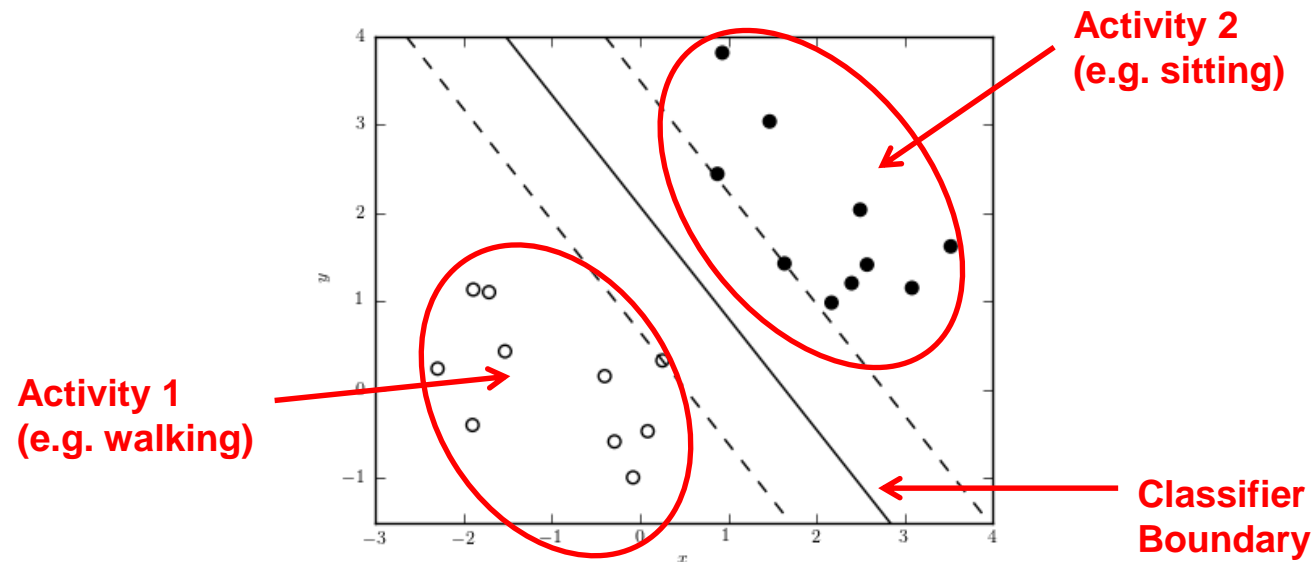
Calculate  
many  
different  
features

- Average[3]: Average acceleration (for each axis)
- Standard Deviation[3]: Standard deviation (for each axis)
- Average Absolute Difference[3]: Average absolute difference between the value of each of the 200 readings within the ED and the mean value over those 200 values (for each axis)
- Average Resultant Acceleration[1]: Average of the square roots of the sum of the values of each axis squared  $\sqrt{(x_i^2 + y_i^2 + z_i^2)}$  over the ED
- Time Between Peaks[3]: Time in milliseconds between peaks in the sinusoidal waves associated with most activities (for each axis)
- Binned Distribution[30]: We determine the range of values for each axis (maximum – minimum), divide this range into 10 equal sized bins, and then record what fraction of the 200 values fell within each of the bins.



## Step 5: Train classifier

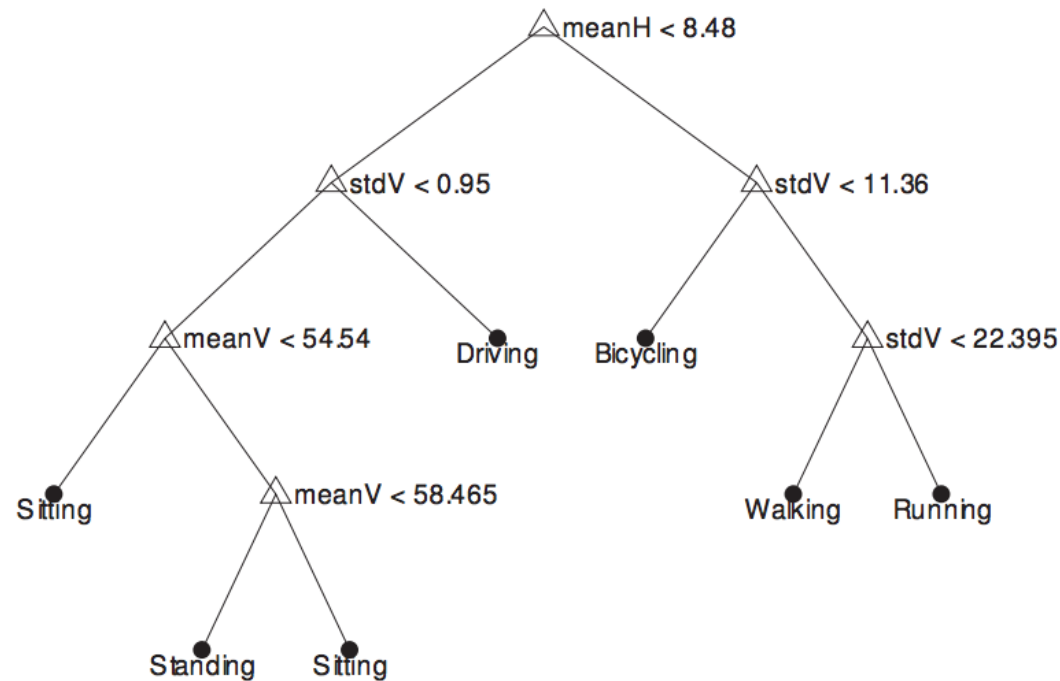
- Feature values are just numbers
- Different feature values for different activities
- **Training classifier:** figures out feature values corresponding to each activity
- Scikit-Learn, MATLAB already programmed with different classification algorithms (SVM, Naïve Bayes, Random Forest, J48, logistic regression, SMO, etc)
- Try different classification algorithms, compare accuracy
- SVM example






## Step 5: Train classifier

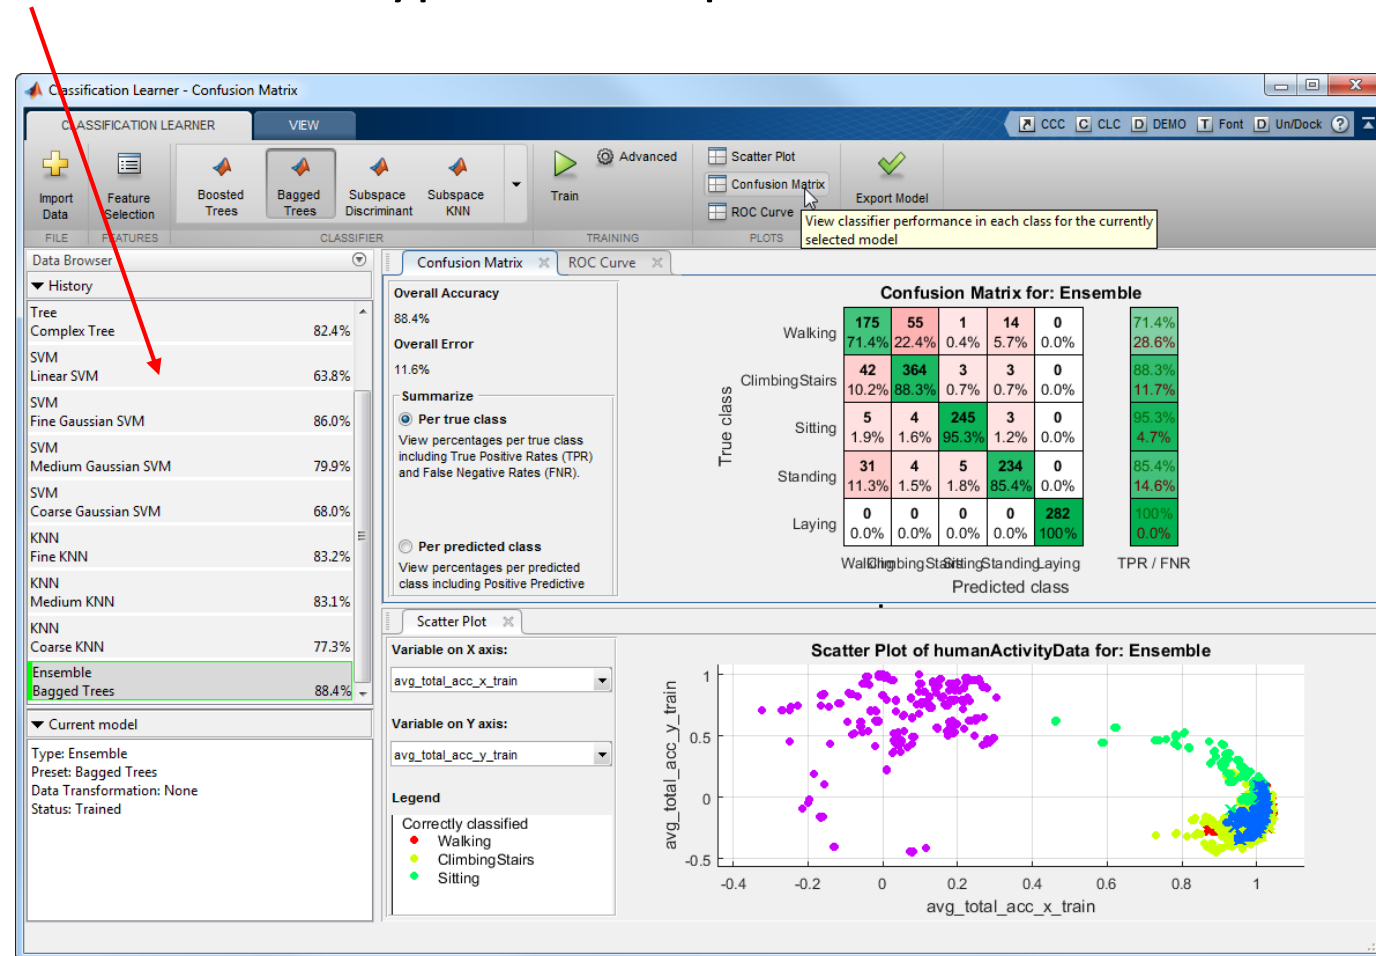
- Typically split data: E.g. 80% for training classifier, 20% for testing
- **Example:** Decision Tree Classifier
- **Training:** Learns thresholds for feature values, which separate the classes
- **Testing:** How well trained model guesses labels (e.g. activity) of subjects in the test set (new examples)





## Step 5: MATLAB Classification Learner App

- Import accelerometer data into MATLAB
- Click and select Classifier types to compare





## Step 5: Train classifier

### Compare Accuracy of Classifier Algorithms

- Scikit-Learn, MATLAB also reports accuracy of each classifier type
- **Accuracy:** Percentage of test cases that classifier guessed correctly

Table 2: Accuracies of Activity Recognition


	% of Records Correctly Predicted			
	J48	Logistic Regression	Multilayer Perceptron	Straw Man
Walking	89.9	<u>93.6</u>	91.7	37.2
Jogging	96.5	98.0	<u>98.3</u>	29.2
Upstairs	59.3	27.5	<u>61.5</u>	12.2
Downstairs	<u>55.5</u>	12.3	44.3	10.0
Sitting	<u>95.7</u>	92.2	95.0	6.4
Standing	<u>93.3</u>	87.0	91.9	5.0
Overall	85.1	78.1	<u>91.7</u>	37.2

Compare, pick most accurate classification algorithm



## Step 6: Deploy Classifier

- Export classification model (most accurate classifier type + data threshold values)
- Classifies new data live!
- Many options to deploy best classifier
- E.g. Program HTTP server, receives data, classifies, sends back results.
- In app write Android code to
  - Gather accelerometer data, segment, extract feature, send features to server



**New accelerometer  
Sample in real time**



Send data  
to server



Get back,  
display results



HTTP web server



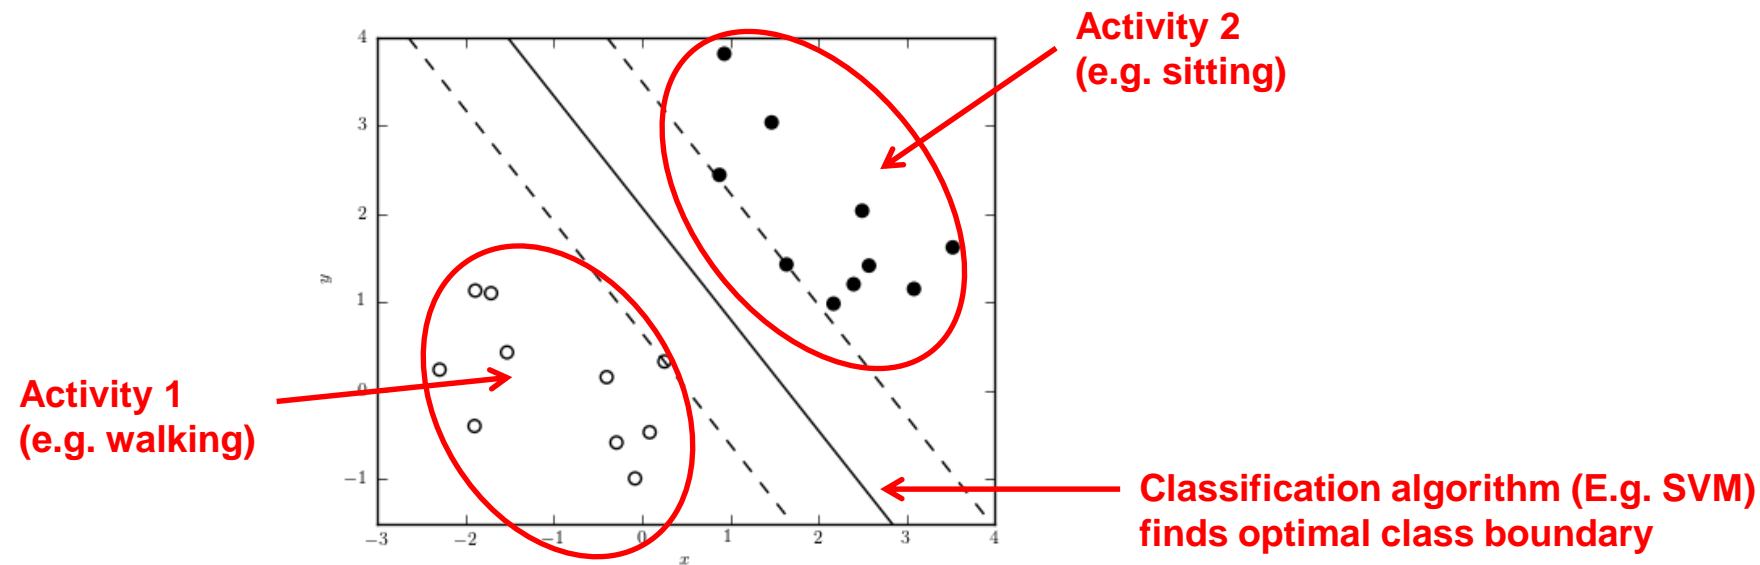
# **Machine Learning Classification Algorithm: Support Vector Machine (SVM)**





# Scalable Vector Machines (SVM)


- Popular machine learning classification algorithm
- **Goal:** Determine optimal boundary between data points corresponding to different classes
- E.g. boundary between data belonging to different activities





# SVM: Delineating Boundaries

- Multiple optimal boundaries exist




**Figure 2. Multiple ways to separate two groups.**



# SVM: Support Vectors

- SVM steps:
  1. Finds **support vectors** in each group: peripheral data points in group 1 that are closest to points in group 2
  2. Find **optimal boundary** between support vectors of both groups
- SVM computationally efficient since it relatively few data points (support vectors)

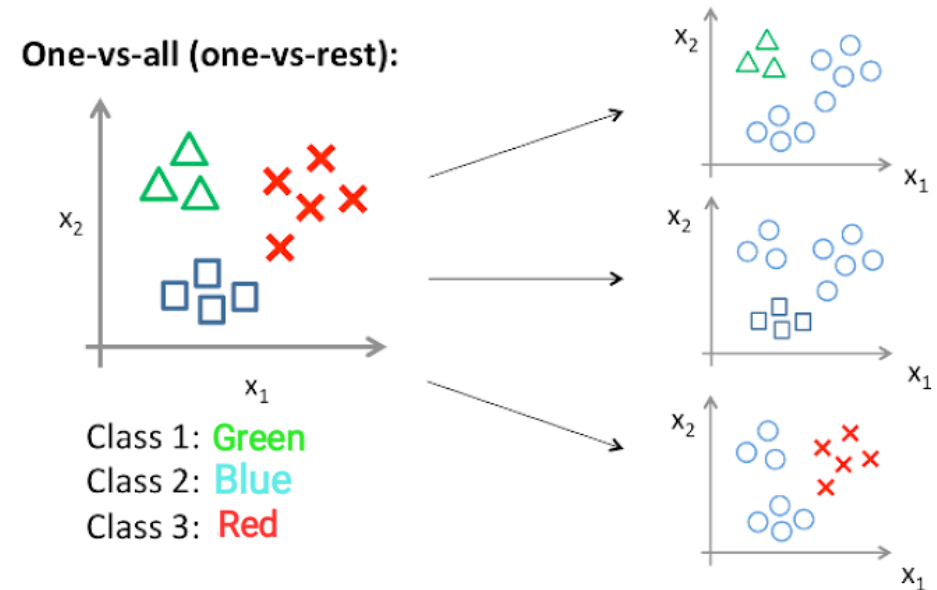



**Figure 3. Optimal boundary is located in the middle of peripheral data points from opposing groups.**




# SVM Limitations

- **Inaccurate for small datasets:** fewer points, less likely to find good support vectors
- **Classifying multiple groups:**
  - SVM classifies 2 groups at a time.
  - Multi-group SVM: Multiple groups handled using multiple 2-group classifications
  - On each iteration, classify 1 group from the rest






# ***k*-Nearest Neighbors Classifier**



# K-Nearest Neighbors

- Assign each point same class as majority of its  $k$  nearest neighbors
- E.g. if  $k = 5$  (below), unknown point (?) classified as red
- Why? 4 red neighbors, 1 black




**Figure 1. The center data point would be classified as red by a majority vote from its five nearest neighbors.**

- $k$  is the number of neighbors to consider for voting

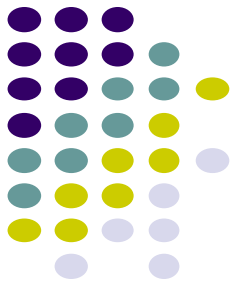


# K-Nearest Neighbors

- $k$  is a parameter, which affects accuracy
  - $k$  too small, algorithm considers only immediate neighbors => overfitting
  - $k$  too large, tries to fit data points too far, not relevant => underfit
- **Overfitting:** algorithm determines boundaries that fits specific dataset but boundary may not hold for new data points



**Figure 2. Comparison of model fit using varying values of  $k$ . Points in the black region are predicted to be white wines, while those in the red region are predicted to be red wines.**



# Context Sensing









## Recall: Ubicomp Senses User's Context

- Context?
  - *Human*: motion, mood, identity, gesture
  - *Environment*: temperature, sound, humidity, location
  - *Computing Resources*: Hard disk space, memory, bandwidth
  - *Ubicomp example*:
    - *Assistant senses*: Temperature outside is 10F (environment sensing) + Human plans to go work (schedule)
    - *Ubicomp assistant advises*: Dress warm!
- Sensed **environment + Human + Computer resources** = **Context**
- *Context-Aware* applications adapt their behavior to context

# Context Sensing

- Activity Recognition typically uses data from 2 sensors: accelerometer and gyroscope
- **User context recognition:** Use machine learning to analyze combined data from multiple sensors (all smartphone sensors?)






# Regression



# Linear Regression

- Strongest predictors of home prices are:
  1. Number of rooms in house
  2. Number of low income neighbors in area
- Linear Regression:
  1. Plot these variables for actual example homes
  2. Fit straight line of best fit
  3. Can use this best fit line to guess price of new homes




**Figure 1. House price against number of rooms.**





# Different Types of Regression

- Different regression functions to fit data to
  - Linear
  - Polynomial
  - Decision tree
  - Etc
- Determine which function has best fit, lowest error (difference)



Polynomial






# Deep Learning



# Deep Learning/Neural Networks


- Machine learning models described so far are traditional/classic ML
  - Involves converting raw data to features (extraction)
- Newer approaches use neural networks or deep learning
  - Learns directly from features or raw data (no need for feature extraction)
- **Neural Networks (NN):** Network of nodes, connectivity weights learned from data
- Learns from data, best weights of edges to classify inputs (x) into outputs y





# Deep Learning/Neural Networks

- NN generally more accurate if adequate data is available
- Requires lots of computational power to train
- NN first outperformed traditional ML on image classification in 2012 (AlexNet)
- For most tasks today (image, speech, text, sensor, etc.) NN solution is most accurate

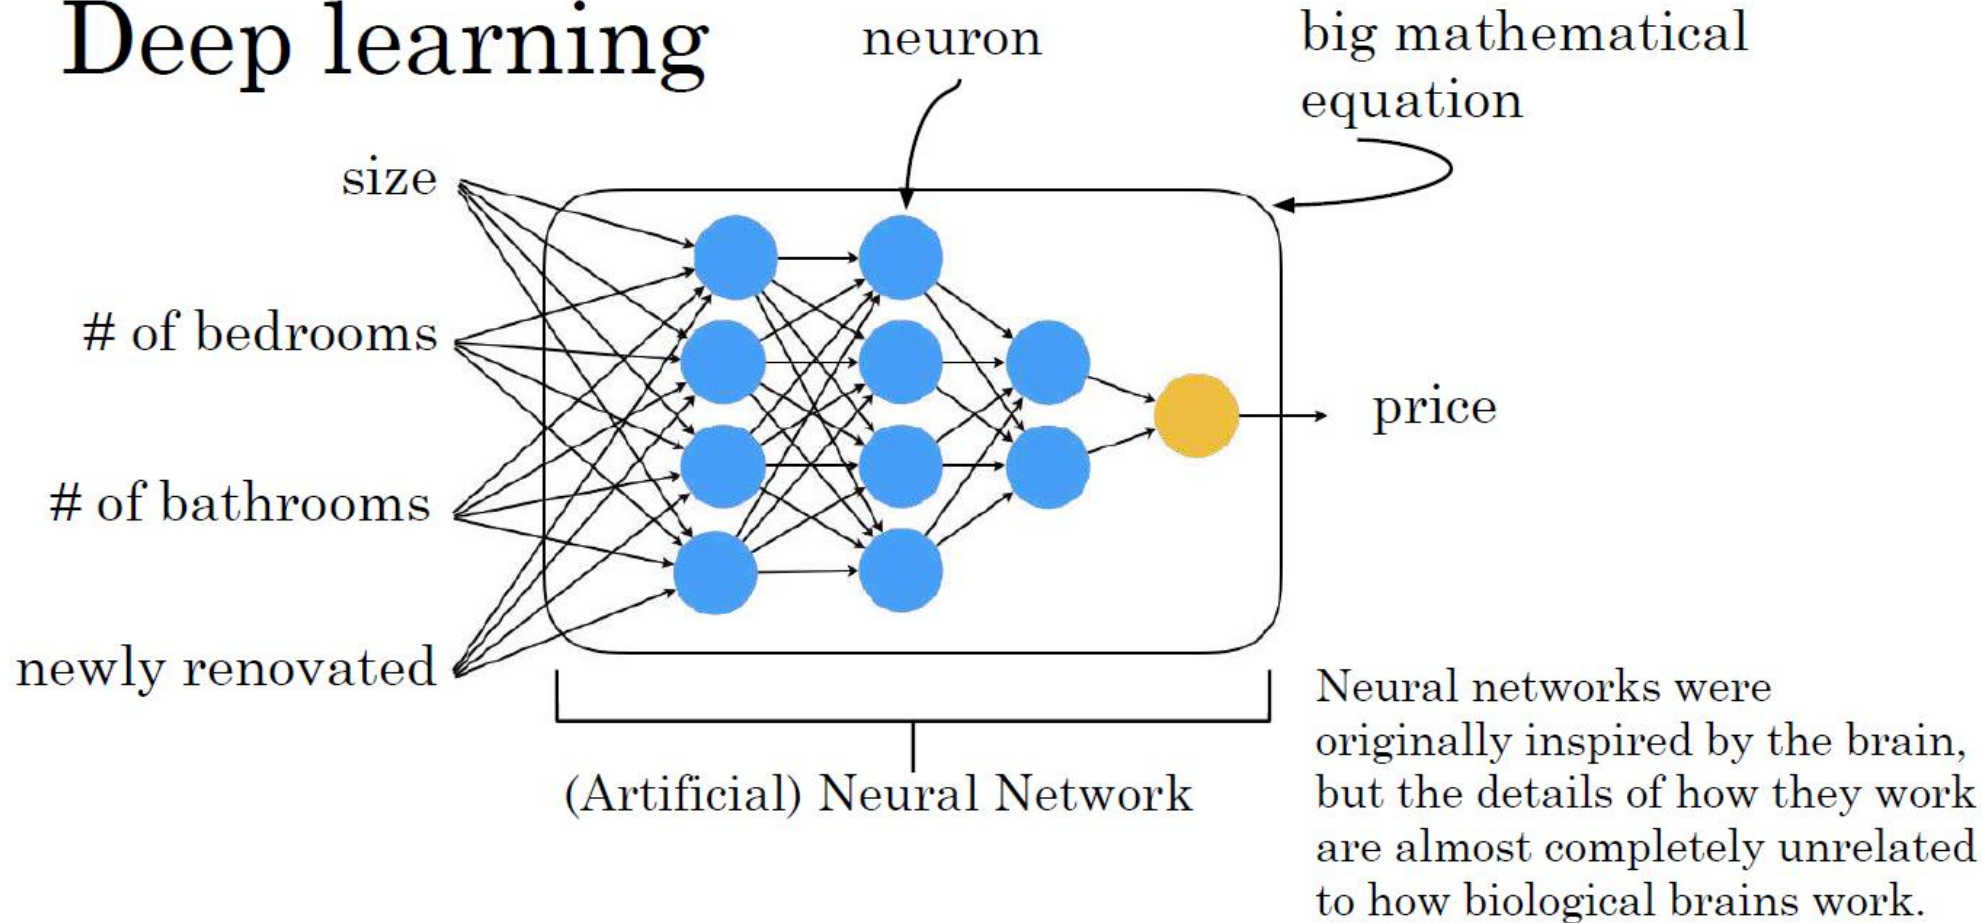




Can conceptualize as NN generates a curve/fitting function to fit data (massive equation)



# Deep learning

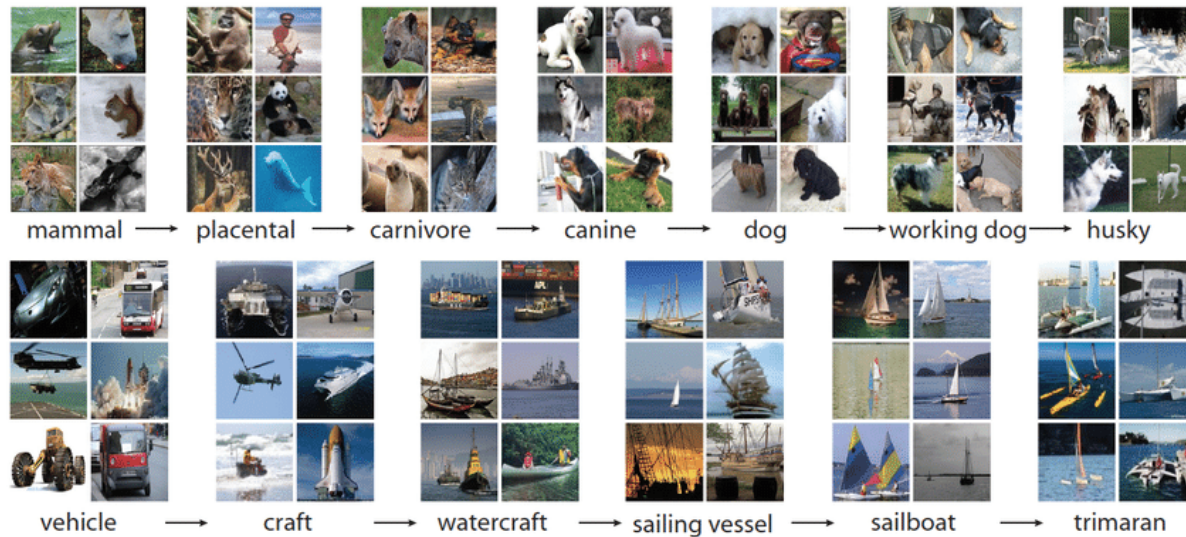


Courtesy: AI for everyone by deeplearning.ai

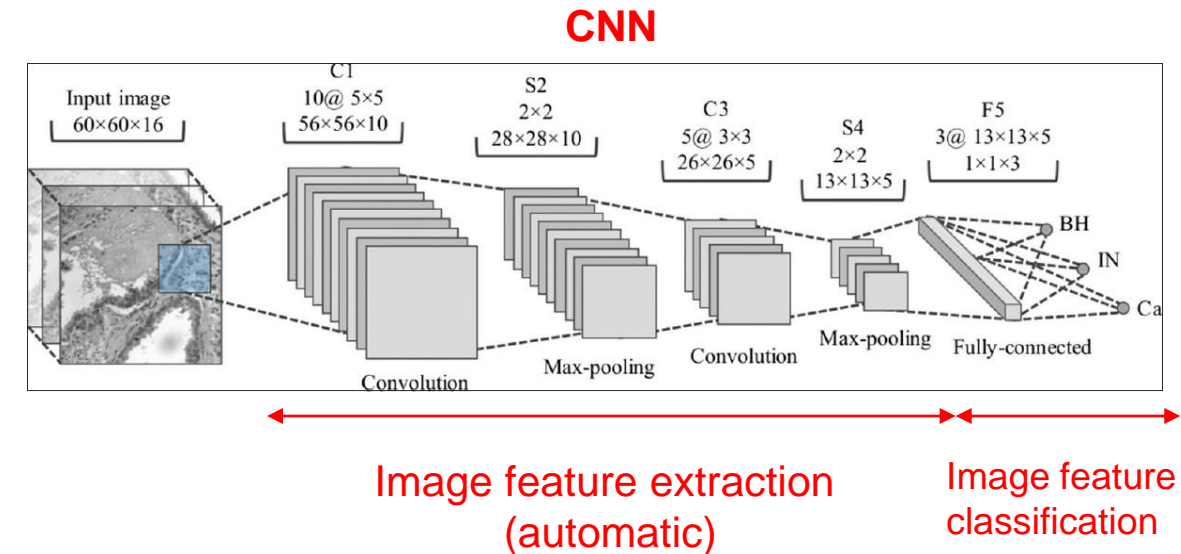


# Convolutional Neural Networks (CNNs)

- Different types of neural networks good for different things
- Convolutional Neural Networks good for classifying images
- E.g. Is there a cat in an input picture?
- **ImageNet**: Popular image dataset, 14 million images in 1000 classes (dogs, sailboat, etc)



ImageNet





# Recurrent Neural Networks (RNNs)

- Good at classifying sequential data
- E.g. Speech translation: sequence of words
- E.g. translate german sentence to English

English: he loves soccer .

Gegman: <start> er ist ein grosser fussballfreund . <end>

Tokenized german: [1, 14, 6, 19, 571, 4200, 3, 2]

English: school's out .

Gegman: <start> die schule ist vorbei . <end>

Tokenized german: [1, 26, 304, 6, 300, 3, 2]

English: i like cookies .

Gegman: <start> ich esse gerne kekse . <end>

Tokenized german: [1, 4, 261, 149, 1656, 3, 2]

English: tom is early .


Gegman: <start> tom ist frueh dran . <end>

Tokenized german: [1, 5, 6, 391, 338, 3, 2]

English: where's tom from ?

Gegman: <start> woher kommt tom ? <end>

Tokenized german: [1, 971, 120, 5, 7, 2]




**RNN**

# Programming/Mobile Support for Neural Networks

<https://developer.android.com/ndk/guides/neuralnetworks>

- Python libraries for neural networks/deep learning, train models in few lines of code
  - Keras
  - PyTorch
  - ScikitLearn
- Training neural networks on Smartphone still tough, currently only testing
- From Android 8.1: Android Neural Networks API (NNAPI) allows inference (test) of pre-trained neural networks on smartphone
  - Supports several machine learning frameworks (e.g. Tensorflow lite)



# References



- Jennifer R. Kwapisz, Gary M. Weiss, and Samuel A. Moore, Activity recognition using cell phone accelerometers, SIGKDD Explor. Newsl. 12, 2 (March 2011), 74-82.
- Deepak Ganesan, Activity Recognition, Physiological Sensing Class, UMASS Amherst